Exam

June 18, 2024

1 Exam 7th of June 2024, 8.00-13.00 for the course 1MS041 (Introduction to Data Science / Introduktion till dataanalys)

1.1 Instructions:

1. Complete the problems by following instructions.
2. When done, submit this file with your solutions saved, following the instruction sheet.

This exam has 3 problems for a total of 40 points, to pass you need 20 points. The bonus will be added to the score of the exam and rounded afterwards.

1.2 Some general hints and information:

- Try to answer all questions even if you are uncertain.
- Comment your code, so that if you get the wrong answer I can understand how you thought this can give you some points even though the code does not run.
- Follow the instruction sheet rigorously.
- This exam is partially autograded, but your code and your free text answers are manually graded anonymously.
- If there are any questions, please ask the exam guards, they will escalate it to me if necessary.

1.3 Tips for free text answers

- Be VERY clear with your reasoning, there should be zero ambiguity in what you are referring to.
- If you want to include math, you can write LaTeX in the Markdown cells, for instance `$f(x)=x^2$` will be rendered as $f(x) = x^2$ and `$$f(x) = x^2$$` will become an equation line, as follows
$$f(x) = x^2$$
Another example is `$$f_{Y \mid X}(y,x) = P(Y = y \mid X = x) = \exp(\alpha \cdot x + \beta)$$` which renders as
$$f_{Y|X}(y,x) = P(Y = y \mid X = x) = \exp(\alpha \cdot x + \beta)$$

1.4 Finally some rules:

- You may not communicate with others during the exam, for example:
  - You cannot ask for help in Stack-Overflow or other such help forums during the Exam.
  - You may not communicate with AI's, for instance ChatGPT.
  - Your on-line and off-line activity is being monitored according to the examination rules.

1.5 Good luck!

```
[ ]: # Insert your anonymous exam ID as a string in the variable below
     examID="XXX"
```

1.6 Exam vB, PROBLEM 1

Maximum Points = 14

In this problem you will do rejection sampling from complicated distributions, you will also be using your samples to compute certain integrals, a method known as Monte Carlo integration: (Keep in mind that choosing a good sampling distribution is often key to avoid too much rejection)

1. [4p] Fill in the remaining part of the function `problem1_rejection` in order to produce samples from the below density using **rejection** sampling:

$$f[x] = Cx^{0.2}(1-x)^{1.3}$$

for $0 \le x \le 1$, where C is a value such that f above is a density (i.e. integrates to one). Hint: you do not need to know the value of C to perform rejection sampling.

2. [2p] Produce 100000 samples (**use fewer if it takes too long**) and put the answer in `problem1_samples` from the above distribution and plot the histogram.

3. [2p] Define X as a random variable with the density given in part 1. Denote $Y = \sin(10X)$ and use the above 100000 samples to estimate

$$\mathbb{E}[Y]$$

and store the result in `problem1_expectation`.

4. [2p] Use Hoeffdings inequality to produce a 95% confidence interval of the expectation above and store the result as a tuple in the variable `problem1_interval`

5. [4p] Can you calculate an approximation of the value of C from part 1 using random samples? Provide a plot of the histogram from part 2 together with the true density as a curve, recall that this requires the value of C. Explain what method you used and what answer you got.

```
[ ]: # Part 1

     def problem1_rejection(n_samples=1):
         # Distribution from part 1
         # write the code in this function to produce samples from the distribution␣
     ↪in the assignment

         # Return a numpy array of length n_samples
         return XXX
```

```
[ ]: # Part 2

     problem1_samples = XXX
```

```
# Part 3

problem1_expectation = XXX
```

```
# Part 4

problem1_interval = [XXX,XXX]
```

```
# Part 5

problem1_C = XXX
```

```
# Part 5

# Write your code to produce the plot here

#XXXXXXX
```

2 Part 5

Double click this cell and directly edit below to answer part 5

2.0.1 Begin explanation

Bla di bla

2.0.2 End explanation

Local Test for Exam vB, PROBLEM 1   Evaluate cell below to make sure your answer is valid. You **should not** modify anything in the cell below when evaluating it to do a local test of your solution. You may need to include and evaluate code snippets from lecture notebooks in cells above to make the local test work correctly sometimes (see error messages for clues). This is meant to help you become efficient at recalling materials covered in lectures that relate to this problem. Such local tests will generally not be available in the exam.

```python
# This cell is just to check that you got the correct formats of your answer
import numpy as np
try:
    assert(isinstance(problem1_rejection(10), np.ndarray))
except:
    print("Try again. You should return a numpy array from problem1_rejection")
else:
    print("Good, your problem1_rejection returns a numpy array")

try:
    assert(isinstance(problem1_samples, np.ndarray))
```

```
except:
    print("Try again. your problem1_samples is not a numpy array")
else:
    print("Good, your problem1_samples is a numpy array")

try:
    assert(isinstance(problem1_expectation, float))
except:
    print("Try again. your problem1_expectation is not a float")
else:
    print("Good, your problem1_expectation is a float")

try:
    assert(isinstance(problem1_interval, list) or isinstance(problem1_interval,␣
↪tuple)) , "problem1_interval not a tuple or list"
    assert(len(problem1_interval) == 2) , "problem1_interval does not have␣
↪length 2, it should have a lower bound and an upper bound"
except Exception as e:
    print(e)
else:
    print("Good, your problem1_interval is a tuple or list of length 2")

try:
    assert(isinstance(problem1_rejection_2(10), np.ndarray))
except:
    print("Try again. You should return a numpy array from problem1_rejection_2")
else:
    print("Good, your problem1_rejection_2 returns a numpy array")
```

2.1 Exam vB, PROBLEM 2

Maximum Points = 13

Let us build a proportional model ($\mathbb{P}(Y = 1 \mid X) = G(\beta_0 + \beta \cdot X)$ where G is the logistic function) for the spam vs not spam data. Here we assume that the features are presence vs not presence of a word, let X_1, X_2, X_3 denote the presence (1) or absence (0) of the words ("$free$","$prize$","win").

1. [2p] Load the file `data/spam.csv` and create two numpy arrays, `problem2_X` which has shape (n_emails,3) where each feature in `problem2_X` corresponds to X_1, X_2, X_3 from above, `problem2_Y` which has shape **(n_emails,)** and consists of a 1 if the email is spam and 0 if it is not. Split this data into a train-calibration-test sets where we have the split 40%, 20%, 40%, put this data in the designated variables in the code cell.

2. [4p] Follow the calculation from the lecture notes where we derive the logistic regression and implement the final loss function inside the class `ProportionalSpam`. You can use the `Test` cell to check that it gives the correct value for a test-point.

3. [4p] Train the model `problem2_ps` on the training data. The goal is to calibrate the prob-

4

abilities output from the model. Start by creating a new variable `problem2_X_pred` (shape `(n_samples,1)`) which consists of the predictions of `problem2_ps` on the calibration dataset. Then train a calibration model using `sklearn.tree.DecisionTreeRegressor`, store this trained model in `problem2_calibrator`.

4. [3p] Use the trained model `problem2_ps` and the calibrator `problem2_calibrator` to make final predictions on the testing data, store the prediction in `problem2_final_predictions`. Compute the $0-1$ test-loss and store it in `problem2_01_loss` and provide a 99% confidence interval of it, store this in the variable `problem2_interval`, this should again be a tuple as in **problem1**.

```python
# Part 1
problem2_X = XXX
problem2_Y = XXX

problem2_X_train = XXX
problem2_X_calib = XXX
problem2_X_test = XXX

problem2_Y_train = XXX
problem2_Y_calib = XXX
problem2_Y_test = XXX


print(problem2_X_train.shape,problem2_X_calib.shape,problem2_X_test.
 ↪shape,problem2_Y_train.shape,problem2_Y_calib.shape,problem2_Y_test.shape)
```

```python
# Part 2

class ProportionalSpam(object):
    def __init__(self):
        self.coeffs = None
        self.result = None

    # define the objective/cost/loss function we want to minimise
    def loss(self,X,Y,coeffs):

        return XXX

    def fit(self,X,Y):
        import numpy as np
        from scipy import optimize

        #Use the f above together with an optimization method from scipy
        #to find the coefficients of the model
        opt_loss = lambda coeffs: self.loss(X,Y,coeffs)
        initial_arguments = np.zeros(shape=X.shape[1]+1)
        self.result = optimize.minimize(opt_loss, initial_arguments,method='cg')
        self.coeffs = self.result.x
```

```
    def predict(self,X):
        #Use the trained model to predict Y
        if (self.coeffs is not None):
            G = lambda x: np.exp(x)/(1+np.exp(x))
            return np.round(10*G(np.dot(X,self.coeffs[1:])+self.coeffs[0]))/10 #␣
↪This rounding is to help you with the calibration
```

```
[ ]: # Part 3

     problem2_ps = XXX

     problem2_X_pred = XXX

     problem2_calibrator = XXX
```

```
[ ]: # Part 4

     # These are the predicted probabilities
     problem2_final_predictions = XXX


     # In order to compute this loss we first need to convert the predicted␣
     ↪probabilities to a decision
     # recall the Bayes classifier?
     problem2_01_loss = XXX

     # Recall the interval is given as a tuple (a,b) or a list [a,b]
     problem2_interval = XXX
```

Local Test for Exam vB, PROBLEM 2   Evaluate cell below to make sure your answer is valid. You **should not** modify anything in the cell below when evaluating it to do a local test of your solution. You may need to include and evaluate code snippets from lecture notebooks in cells above to make the local test work correctly sometimes (see error messages for clues). This is meant to help you become efficient at recalling materials covered in lectures that relate to this problem. Such local tests will generally not be available in the exam.

```
[ ]: try:
         import numpy as np
         test_instance = ProportionalSpam()
         test_loss = test_instance.loss(np.array([[1,0,1],[0,1,1]]),np.
     ↪array([1,0]),np.array([1.2,0.4,0.3,0.9]))
         assert (np.abs(test_loss-1.2828629432232497) < 1e-6)
         print("Your loss was correct for a test point")
     except:
```

```
        print("Your loss was not correct on a test point")
```

2.2 Exam vB, PROBLEM 3

Maximum Points = 13

Consider the following four Markov chains, answer each question for all chains:

Markov chain A Markov chain B Markov chain C Markov chain D

1. [2p] What is the transition matrix?
2. [2p] Is the Markov chain irreducible?
3. [3p] Is the Markov chain aperiodic? What is the period for each state? Hint: Recall our definition of period; Let $\mathbb{T} := \{t \in \mathbb{N} : P^t(x, x) > 0\}$ and the greatest common divisor of $\mathbb{T}$ is the period.
4. [3p] Does the Markov chain have a stationary distribution, and if so, what is it?
5. [3p] Is the Markov chain reversible?

```
[ ]:  # PART 1

      #-----------------------TRANSITION MATRIX -------------------------------
      # Answer each one by supplying the transition matrix as a numpy array
      # of shape (n_states,n_states), where state (A,B,...) corresponds to index (0,1,.
       ↪..)

      problem3_A    = XXX
      problem3_B    = XXX
      problem3_C    = XXX
      problem3_D    = XXX
```

```
[ ]:  # PART 2
      #-----------------------REDUCIBLE ----------------------------------
      # Answer each one with a True or False

      problem3_A_irreducible = XXX
      problem3_B_irreducible = XXX
      problem3_C_irreducible = XXX
      problem3_D_irreducible = XXX
```

```python
# PART 3
#----------------------APERIODIC-----------------------------
# Answer each one with a True or False

problem3_A_is_aperiodic = XXX
problem3_B_is_aperiodic = XXX
problem3_C_is_aperiodic = XXX
problem3_D_is_aperiodic = XXX

# Answer the following with the period of the states as a numpy array
# of shape (n_states,)

problem3_A_periods = XXX
problem3_B_periods = XXX
problem3_C_periods = XXX
problem3_D_periods = XXX
```

```python
# PART 4
#-----------------------STATIONARY DISTRIBUTION-----------------
# Answer each one with a True or False

problem3_A_has_stationary = XXX
problem3_B_has_stationary = XXX
problem3_C_has_stationary = XXX
problem3_D_has_stationary = XXX

# Answer the following with the stationary distribution as a numpy array of
 ↪shape (n_states,)
# if the Markov chain has a stationary distribution otherwise answer with False

problem3_A_stationary_dist = XXX
problem3_B_stationary_dist = XXX
problem3_C_stationary_dist = XXX
problem3_D_stationary_dist = XXX
```

```python
# PART 5
#-----------------------REVERSIBLE-----------------
# Answer each one with a True or False

problem3_A_is_reversible = XXX
problem3_B_is_reversible = XXX
problem3_C_is_reversible = XXX
problem3_D_is_reversible = XXX
```